

# MTH 221

## Fundamentals of Machine Learning

Batuhan Bardak

### **Lecture 2:** Basic concepts of ML & ML by example (k-NN)

**Date:** 10.10.2023

# Plan for today

- Recap on ML Basics
- KNN
- Curse of Dimensionality
- Hands on examples with Python (Demo of KNN & Curse of Dimensionality)
- Feature Scaling & Evaluation Metrics

# Grading

- Midterm: %20 (closed book, no notes, no cheat sheet)
- Final: %30 (closed book, no notes, no cheat sheet)
- Homeworks: %20 (with Python)
- Course Project: %30 (done in groups of 2)

**Note:** If you submit your project final reports to the IEEE conference, 20 points will be added to your project grade.


**Note2:** Extra %5 points for active participation in the class.



# Recap on ML Basics

Dataset:

$$\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}, x_i \in \mathbb{R}^d, y_i \in \mathcal{C} \text{ (e.g., } \mathcal{C} = \{-1, 1\}\text{)}, (x_i, y_i) \sim \mathcal{P}$$



Hypothesis:

$$h : \mathbb{R}^d \mapsto \mathcal{C}$$

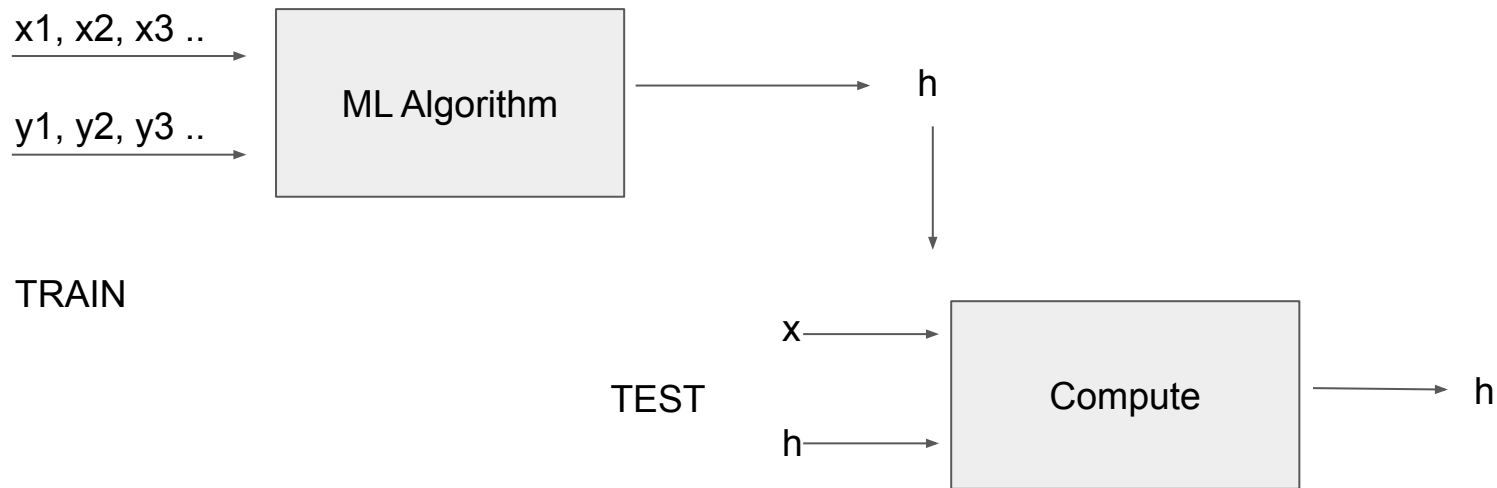
i.e., a neural network-based classifier that maps image to label of cat or dog

Hypothesis class

$$\mathcal{H} = \{h\}$$

i.e., a large family of NNs with different parameters


# Training



# Summary

1. Given a task and a dataset

$$\mathcal{D} = \{x_i, y_i\}, x_i, y_i \sim \mathcal{P}$$



2. Design hypothesis class  $\mathcal{H}$  and loss function  $\ell$  (encodes inductive bias)



3. Train:  $\hat{h} = \arg \min_{h \in \mathcal{H}} \sum_{(x,y \in \mathcal{D})} \ell(h, x, y)$





Output:  $\hat{h}$  that has small generalization error

$$\mathbb{E}_{x,y \sim \mathcal{P}}[\ell(\hat{h}, x, y)]$$


Often repeated many times using  $\mathcal{D}_{VA}$  / cross validation



# Train / Test Split

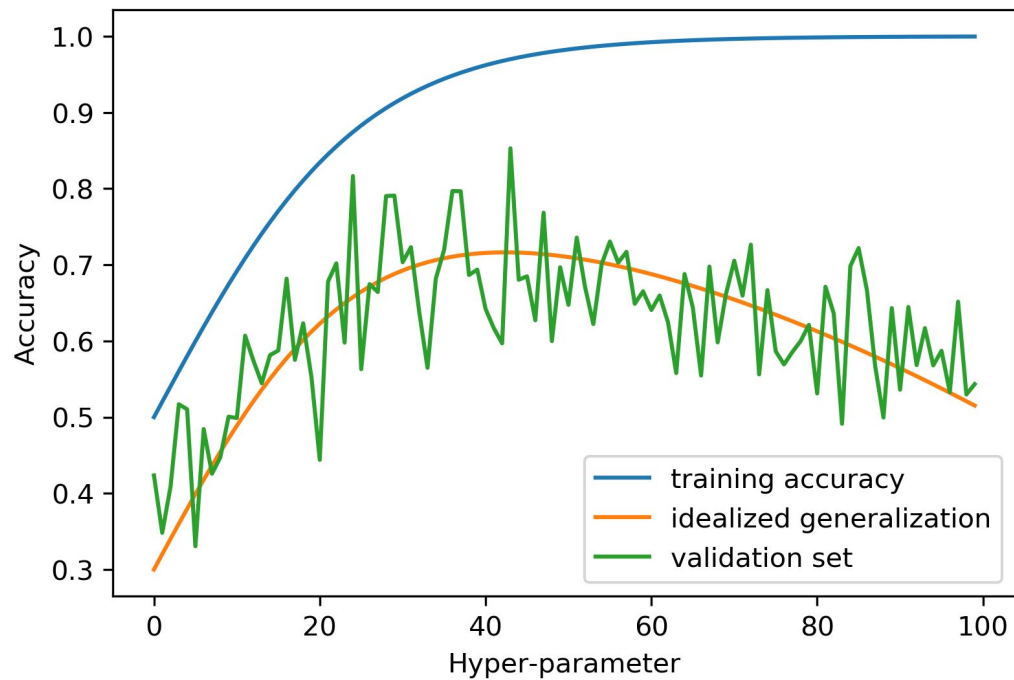


# Train / Test Split







# Train / Test Split




# Overfitting / Underfitting




# Overfitting / Underfitting



# Overfitting / Underfitting



# Cross-Validation



# Recap on ML Basics

- T/F: A hypothesis that achieves zero training error is always good
- T/F: zero-one loss is a good loss function for regression
- T/F: We can use validation dataset to check if our model overfits

# K Nearest Neighbors

# The K-NN Algorithms

**Input:** classification training **dataset**  $\{x_i, y_i\}_{i=1}^n$ , and parameter  $K \in \mathbb{N}^+$ ,  
and a **distance metric**  $d(x, x')$  (e.g.,  $\|x - x'\|_2$  euclidean distance)

## K-NN Algorithm:

Store all training data

For any test point  $x$  :

Find its top K nearest neighbors (under metric  $d$ )

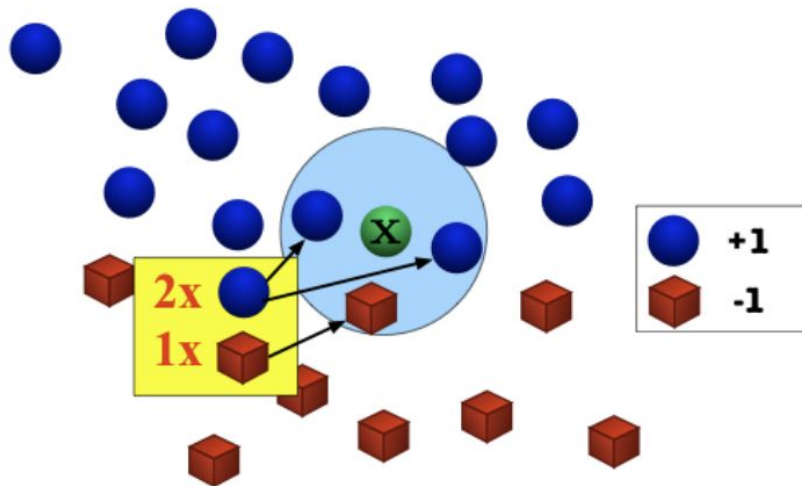
Return the most common label among these K neighbors

(If for regression, return the average value of the K neighbors)



# The K-NN Algorithms

Example: 3-NN for binary classification using Euclidean distance



# The K-NN Algorithms

1. We believe our metric  $d$  captures similarities between examples:

Examples that are close to each other share similar labels

**Another example: Manhattan distance ( $\ell_1$ )**

$$d(x, x') = \sum_{j=1}^d |x[j] - x'[j]|$$

# Distance Metrics

- Euclidean Distance


$$d(p, q) = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2}.$$

- Manhattan Distance

$$d_1(\mathbf{p}, \mathbf{q}) = \|\mathbf{p} - \mathbf{q}\|_1 = \sum_{i=1}^n |p_i - q_i|,$$

- Minkowski Distance

$$D(X, Y) = \left( \sum_{i=1}^n |x_i - y_i|^p \right)^{\frac{1}{p}}.$$



scikit-learn

**p : int, default=2**

Power parameter for the Minkowski metric. When  $p = 1$ , this is equivalent to using `manhattan_distance` (l1), and `euclidean_distance` (l2) for  $p = 2$ . For arbitrary  $p$ , `minkowski_distance` (l\_p) is used.

# The choice of $K$

## 1. What if we set $K$ very large?

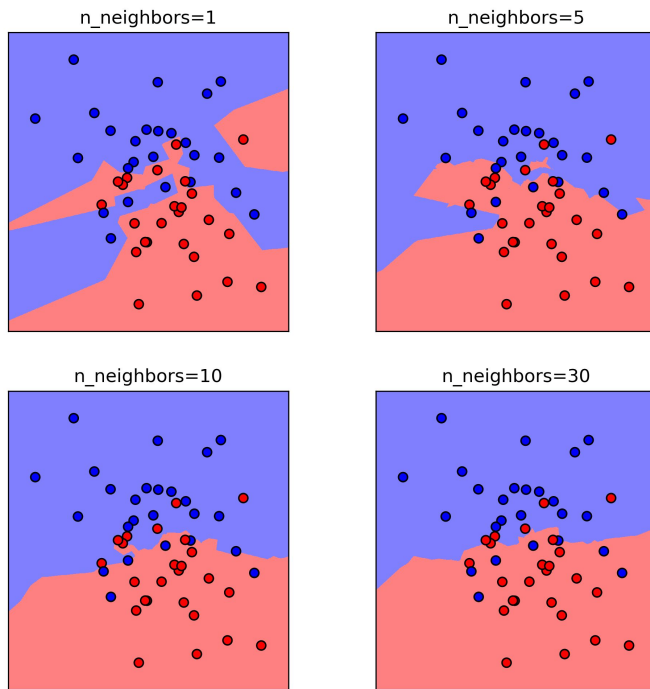
Top  $K$ -neighbors will include examples that are very far away...

## 2. What if we set $K$ very small ( $K=1$ )?

label has noise (easily overfit to the noise)

(What about the training error when  $K = 1$ ?)

# k-Nearest Neighbors



# k-Nearest Neighbors

- Main advantage
  - Simple to understand and implement
  - Immediately adapts as we collect new training data
- Downside
  - The computational complexity for classifying new examples grows linearly with the number of examples in the training dataset in the worst-case scenario
  - We can't discard training examples since no training step is involved. Thus, storage space can become a challenge if we working with large datasets.
  - The algorithm's performance might degrade with high-dimensionality data (aka **Curse of Dimensionality!**)
  - It's sensitive to irrelevant or redundant features since all features contribute equally to the distance calculation (you may use weight to solve this problem!)


# Code Example

- 1-Introduction-to-Python.ipynb
- 2-KNN-Implementation.ipynb
- 3-Find Best K Value for KNN.ipynb
- Check this [link](#) for hyperparameters of kNN.

# Curse of Dimensionality

Key problem: in high dimensional space, points that are drawn from a distribution tend to be far away from each other!

**Example: let us consider uniform distribution over a cube  $[0,1]^d$**



Q: sample  $x$  uniformly, what is the probability that  $x$  is inside the small cube?

A:  $\text{Volume}(\text{small cube}) / \text{volume}([0,1]^d) = l^d$




# Curse of Dimensionality

**Example: let us consider uniform distribution over a cube  $[0,1]^d$**

Now assume we sampled  $n$  points uniform randomly, and we observed  $K$  points fall inside the small cube

So empirically, the probability of sampling a point inside the small cube is roughly  $K/n$




Thus, we have  $l^d \approx \frac{K}{n}$

# Curse of Dimensionality

**Example: let us consider uniform distribution over a cube  $[0,1]^d$**

We have  $l^d \approx \frac{K}{n}$



Q: how large we should set  $l$ , s.t., we will have  $K$  examples (out of  $n$ ) fall inside the small cube?


$$l \approx (K/n)^{1/d} \rightarrow 1, \text{ as } d \rightarrow \infty$$

**Bad news: when  $d \rightarrow \infty$ , the  $K$  nearest neighbors will be all over the place!  
(Cannot trust them, as they are not nearby points anymore!)**

# Can we just increase “n” to avoid this?

**Example: let us consider uniform distribution over a cube  $[0,1]^d$**

We have  $l^d \approx \frac{K}{n}$



Q: to make sure that we have one sample inside a small cube, how large  $n$  needs to be?

Set  $\ell = 0.1$ ,  $K = 1$ , then  $n = 1/(0.1)^d = 10^d$

**Bad news: when  $d \geq 100$ , # of samples needs to be larger than total # of atoms in the universe!**

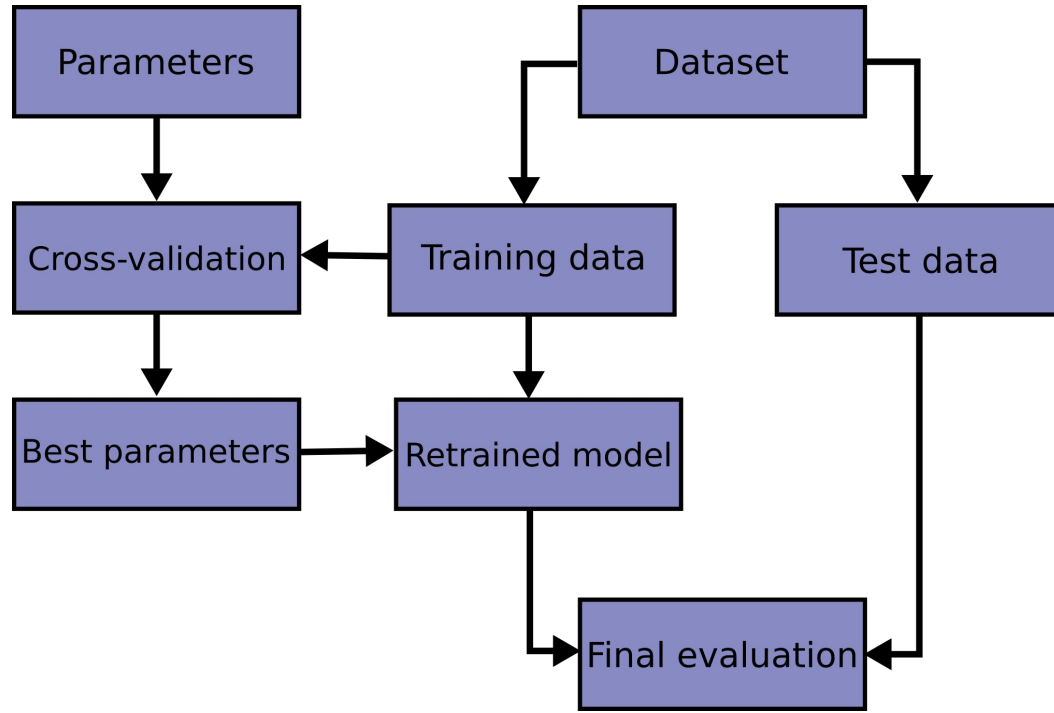
# The curse of dimensionality

- As dimensions increase ..
  - Euclidean distances become less meaningful
  - Uniform distributions become exponentially harder to sample
  - Data becomes more difficult to visualize

# Code Example

- 4-KNN & Curse\_of\_Dimensionality.ipynb

# Pipeline for hyper parameter tuning

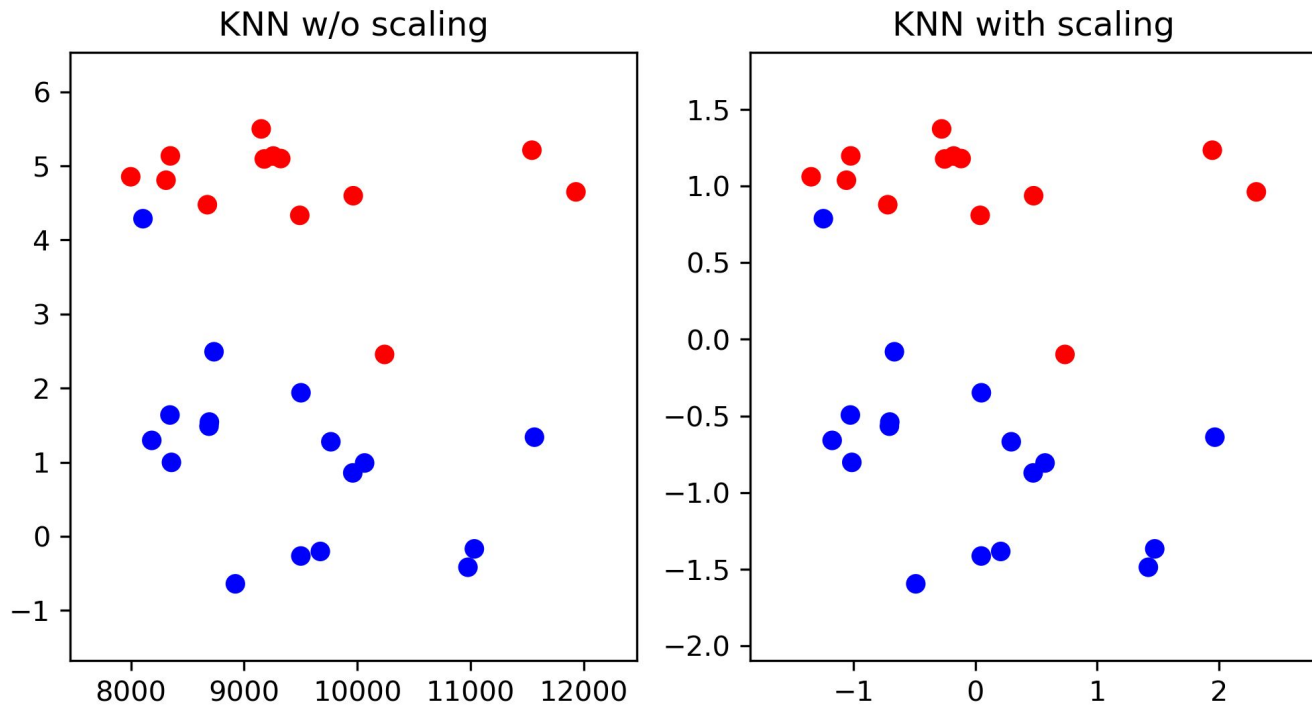


# Further Reading on model evaluation & selection

- Model Evaluation, Model Selection, and Algorithm Selection in Machine Learning
  - <https://arxiv.org/pdf/1811.12808.pdf>
- Cross-validation failure: small sample sizes lead to large error bars
  - <https://hal.inria.fr/hal-01545002/file/paper.pdf>


# Feature Scaling

•






# Feature Scaling



# Different Methods for Scaling Data



# Feature Scaling (Normalization)

- MinMax Scaler

- Transformers features between 0 and 1.

$$\frac{x_i - \min(x)}{\max(x) - \min(x)}$$

- Standard Scaler

- Assumes a normal distribution for data within each feature.
- Scaling makes the distribution centered around 0, with standard deviation of 1 and the mean removed.

$$\frac{x_i - \text{mean}(x)}{sd(x)}$$

- Robust Scaler

- Scale features that are robust to outliers. Similar to MinMax Scaler but it uses the interquartile range (rather than the min-max). The median and scales of the data are removed by this scaling algorithm according to the quantile range.

$$\frac{x_i - Q_1(x)}{Q_3(x) - Q_1(x)}$$

- For comparing methods

- [https://scikit-learn.org/stable/auto\\_examples/preprocessing/plot\\_all\\_scaling.html](https://scikit-learn.org/stable/auto_examples/preprocessing/plot_all_scaling.html)




# Improperly Scaled Data

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scalerfit.transform(X_test)
```



# Code Example

- 5-Feature Scaling.ipynb

# Evaluation Metrics for Binary Classification

actual negative	True Negative	False Positive
actual positive	False Negative	True Positive
	predicted negative	predicted positive

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{TP} + \text{TN} + \text{FP} + \text{FN}}$$

# Problems with Accuracy


Data with 90% negatives:

```
from sklearn.metrics import accuracy_score
for y_pred in [y_pred_1, y_pred_2, y_pred_3]:
    print(accuracy_score(y_true, y_pred))
```

0.9

0.9

0.9



# Evaluation Metrics for Classification

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

Positive Predicted Value (PPV)

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

Sensitivity, coverage, true positive rate.

$$F = 2 \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

Harmonic mean of precision and recall



# Code Example

- 6-Evaluation Metrics.ipynb

# Administrative

- You can do the projects in teams of 2 people.
- Please email the team member(s) till 13.10.2023 and send your project proposal till 20.10.2023.
  - To: [cinarb@metu.edu.tr](mailto:cinarb@metu.edu.tr)
  - Cc: [bbatuhan@metu.edu.tr](mailto:bbatuhan@metu.edu.tr)
- The project proposal should be max one page (min. 300 words) .Your project proposal should describe:
  - What is the problem that you will be investigating? Why is it interesting?
  - What reading will you examine to provide context and background?
  - What data will you use? If you are collecting new data, how will you do it?
  - What method or algorithm are you proposing? If there are existing implementations, will you use them and how? How do you plan to improve or modify such implementations? You don't have to have an exact answer at this point, but you should have a general sense of how you will approach the problem you are working on.
  - How will you evaluate your results? Qualitatively, what kind of results do you expect (e.g. plots or figures)? Quantitatively, what kind of analysis will you use to evaluate and/or compare your results?
- Kaggle (active competition) | Algorithmic Trading | Other

Please submit your proposal as a PDF. **Only one person on your team should** submit. Please have this person add the rest of your team as collaborators as a "Group Submission".

# Administrative

- Expect some questions from the paper below in the midterm exam. (Reading assignment)
  - A Few Useful Things to Know about Machine Learning, *P. Domingos*
    - <https://homes.cs.washington.edu/~pedrod/papers/cacm12.pdf>

Next Class:

Perceptron